

ONNX Training Proposal

AI Frameworks Team

Microsoft

Wei-Sheng Chin

Motivation

- **Training** can be a part of model deployment.
 - Reinforcement learning.
 - Transferring learning.
- User should be able to use whatever training framework/hardware they want to use to train a model.
 - Defining training **locally** and submitting it to a powerful cluster.
- To make AI exchangeable, ONNX should be expressive enough to cover training.

What is training?

- Given a data set, training is a procedure to improve a randomly initialized model by solving an **optimization** problem.
 - The trained model should fit the given data set better than its initial version.
- Here we focus on **gradient**-based methods due to its dominance in this field.

A Common Formulation to Train Neural Network

- Given some data points, we find an (sub-)optimal model by minimizing a quality measure.

$$\min_{w \in R^m} \sum_{i=1}^l L(x_i, w) - (1)$$

- w denotes all trainable model parameters.
- x_i is the i -th data point's feature vector.
- The measure, L , is a loss function, for example, squared distance in regression problem.

A Stochastic Gradient Method to Solve (1)

- *for* $t = 0, \dots,$
 - Randomly pick up $i \in \{1, \dots, l\}$
 - $G \leftarrow \nabla_w L(x_i, w)$
 - $w \leftarrow w - \eta G$
- Some observations
 - A **loss** function, L , we want to minimize.
 - **Gradient** is computed at each iteration based on the latest model w .
 - Gradient need to be **manipulated** before applying it.
 - The model w is updated by an **assignment**.
 - The pattern **repeats** at each iteration.

Fundamental Components to Support Training (# [2038](#))

- Loss function.
- Gradient computation.
- Optimizer for operators.
- Assignment semantic.
- Some spec changes to accommodate these new concepts.

Loss functions (PR #[1939](#))

- We will support popular ones officially.
 - Mean squared error, $(y - \hat{y})^2$.
 - Absolute loss, $|y - \hat{y}|$.
 - Hinge loss, $\max(0, 1 - y\hat{y})$.
 - Logistic loss, $\log(1 - e^{-y\hat{y}})$.
 - Softmax cross entropy loss, $\log(\frac{e^{y_c}}{\sum e^{y_{c'}}})$.
- They should be **composed** as FunctionProto's by existing operators.

Gradient Operator (PR #[2168](#))

- Attributes
 - xs: Names of the n differentiated tensors, for example, $["x_1", "x_2", \dots, "x_n"]$.
 - y: Name of target tensor, for example, "y".
- Inputs
 - Values of the source tensors. The i -th tensor in "Inputs" would be bound to the i -th name in "xs".
- Outputs
 - $[\frac{\partial y}{\partial x_1}, \frac{\partial y}{\partial x_2}, \dots, \frac{\partial y}{\partial x_n}]$.
- All outputs are optional. User can use empty string to indicate unnecessary ones.

Targeted Optimizers

- The number of Pytorch and TF optimizers is **huge**, so it's not realistic to support all of them in the beginning.
- A list of our top priorities, selected by popularity and mathematical properties.
 - Momentum (PR # [1959](#))
 - Adagrad (PR # [1955](#))
 - Adam (PR # [1970](#))

Optimizer's Reference Implementation

- Each PR contains a numpy **reference** implementation.

```
def apply_adagrad(r, t, x, g, h, norm_coefficient, epsilon, decay_factor):  
    # Compute adjusted learning-rate.  
    r_ = r / (1 + t * decay_factor)  
    # Add gradient of regularization term.  
    g_regularized = norm_coefficient * x + g  
    # Update squared accumulated gradient.  
    h_new = h + g_regularized * g_regularized  
    # Compute ADAGRAD's gradient scaling factors  
    h_sqrt = np.sqrt(h_new) + epsilon  
    # Apply ADAGRAD update rule.  
    x_new = x - r_ * g_regularized / h_sqrt  
    return (x_new, h_new)
```

Correctness of Optimizer Operators

- Reference implementation's outputs are compared against both of Pytorch and Tensorflow in each PR.

Mix-precision Training

- Some recent researches show that some float32 parameters can be stored as float16.
- Low-precision numbers means less **computation** and less **memory**.
- ONNX optimizers should allow, for example, weights in float32 and gradient in float16.
 - It's not difficult. We just allows more types in optimizers' input list.

Storing Training Algorithm

- Training algorithm is broken into stages so that **no assignment** happens in each stage.
- One stage is just a computation graph (type: **FunctionProto**).
 - Gradient node + optimizer nodes + loss node + etc.
- Assignment is independently encoded outside stage's actual computation.
- **Sequentially** executing **all** stages is considered as one training iteration.

Storing Training Stages in ModelProto

- message ModelProto {
 ...
 // The parameterized graph that is evaluated to execute the model.
 optional GraphProto graph = 7;
 // The functions can be referenced in the "graph.node" field and
 // "training_info[i].algorithm.node" field.
 repeated FunctionProto function = 9;
 ...
 // Training-specific information. Sequentially executing all stored
 // "TrainingInfoProto"s is one training iteration.
 repeated TrainingInfoProto training_info = 21;
}

Storing Training Stage in TrainingInfoProto

- message TrainingInfoProto {
 // Training-specific initializers such as learning rate.
 repeated TensorProto **initializer** = 1;
 // Algorithm's inputs.
 repeated ValueInfoProto **input** = 2;
 // Selected outputs from "algorithms"
 repeated ValueInfoProto **output** = 3;
 // The training algorithm, accessing **input**, **initializer**, **ModelProto.graph.initializer**,
 ModelProto.graph.function.
 optional FunctionProto **algorithm** = 4;
 // String pair such as (" W_{new} ", " W ") to encode assignment semantic " $W = W_{new}$ ".
 repeated StringStringEntryProto **update_binding** = 7;
}

Summary Message on GitHub

Summary of Training Story in ONNX #2038

 Open

wschin opened this issue on May 21 · 2 comments



wschin commented on May 21 • edited ▼

Member



This issue summarizes PRs related ONNX Training WG's progress.

We have a PR [#2013](#) for defining required protobuf changes to enable training. That PR creates a new protobuf message to store `optimizer`, `loss`, and other useful fields. A graphical example can be found in [ONNX Training Discussion.pptx](#)

For loss, one can see [#1939](#).

For distinguishing behaviors of operators in training and inference phases, see [#1887](#).

For optimizers, we have [#1955](#), [#1970](#), [#1959](#).

For gradient computation, [#2168](#).



1

Thank you!